April 18, 2017
DS-GA 3001-004, Text as Data
Prof Arthur Spirling

Homework 3

This homework must be returned to Patrick Chester's mailbox at 19 West 4th St at **5:00 PM on May 5th, 2017**. Late work will incur penalties of the equivalent of one third of a letter grade per day late.

It must be your own work, and your own work only — you must not copy anyone's work, or allow them to copy yours. This extends to writing code. You may consult with others, but when you write up, you must do so alone.

You must include an `R` code appendix which is clearly commented, such that it your understanding of the problems can be assessed. You must turn in a paper copy: **no electronic copies will be accepted**.

For the following exercises, it is recommended that you use the following packages: `topicmodels`, `lda`, and `stm`.

---

1. Applying `topicmodels` to the news corpus

   (a) To decrease the time it takes to fit a topic model, we will limit our analysis to a subset of the immigration corpus. Create a subset of `data_corpus_immigrationnews` that only contains articles from the 4 news sources with the most documents in the immigration news corpus. Create a table that shows how many documents are associated with each newspaper.

   (b) Create a document term matrix with your new immigration corpus in which punctuation is removed and words are set to lower case. Also, remove a custom set of stopwords `custom_stopwords` - which can be found in the homework folder - that is relevant to this particular data set. Finally, use `quanteda`'s "dfm_trim" to remove words that occur fewer than 30 times or in fewer than 20 documents. Print the remaining number of words and that total number of documents in the DFM.

   (c) Preprocessing decisions can have substantive impacts on the topics created by topic model algorithms. Make a brief (1 paragraph) argument for or against removing rare terms from a dfm on which you plan to fit a topic model.

   (d) Fit a topic model with 30 topics using `LDA()`, with `method = "Gibbs"`. Increase the number of iterations to 3000 to ensure that the model describes the underlying data well and set the seed to 10012 so that you can replicate your results.

(e) Examine the words that contribute the most to each topic (examine the top 10). Use `get_terms()`. Choose names for each of the 5 topics that have the most articles in the corpus associated with them. Explain your choice of labels. Save the top 10 words from all 30 topics, for later use.

(f) Examine the topics that contribute the most to each document, using the code from Recitation_11 to visualize the top two topics per document for the Guardian and the Telegraph with separate graphs for each newspaper. Make sure that the documents are sorted by day of publication (the "day" variable in the `data_corpus_immigrationnews` corpus). Discuss your findings.

(g) Finally, we can find the average contribution of a topic to an article from a particular newspaper, and compare newspapers on particular topics. For each of the 5 topics you've named, see how their prevalence varies among the different newspapers. To do so, estimate the mean contribution of each topic over each newspaper. Discuss your findings.

2. **Topic Stability:** We want to see how stable these topics are, under two different sets of pre-processings

(a) Re-run the model from question 1 with a different seed.

(b) For each topic in the new model, find the topic that is the closest match in the original run in terms of cosine similarity of the topic distribution over words.

(c) Calculate the average number of words in the top ten shared by each matched topic pair.

(d) Now run two more models, but this time, use only 5 topics. Again, find the average number of words in the top ten shared by each matched topic pair. How stable are the models with 5 topics compared to the models with 30 topics?

3. **Topic Models with covariates:** The Structural Topic Model (STM) is designed to incorporate document-level variables into a standard topic model. Since, we have information about both the newspaper and the date of the articles, we can use an STM (from the `stm` package) to model the effects of these covariates directly.

(a) Using only articles from the Guardian and Telegraph, construct a numeric date variable from the "day" variable in the immigration news corpus. Use what preprocessing you believe to be appropriate for this problem.

(b) Fit an STM model where the topic content varies according to this binary variable, and where the prevalence varies according to both this binary variable and the spline of the

date variable you've created. Be sure to use the spectral initialization and set k=0 which will allow the STM function to automatically select a number of topics using the *spectral learning* method. Keep in mind that this function is computationally demanding, so start with the minimum threshold document frequency threshold set to 10; if your computer takes an unreasonably long time to fit the STM model with this threshold, you can raise it to as high as 30.

(c) Identify and name each of the 5 topics that occur in the highest proportion of documents using the following code:[1]

```
plot(fit.stm, type = "summary")
```

(d) Using the visualization commands in the `stm` package, discuss one of these top 5 topics. How does the content vary with the paper discussing that topic? How does the prevalence change over time?

4. **Non-Parametric Scaling - Wordfish:** Recall that the Wordfish algorithm allows us to scale political texts by a latent dimension. We will apply this function to analyze the State of the Union addresses.

(a) First, create a corpus that is the subset of the `data_corpus_SOTU` that contains only speeches that occurred after 1970.

(b) Wordfish requires that we select anchors that lie at the extremes of the latent dimension; in this case, we are looking to estimate the latent left-right ideological dimension. Use Obama's 2012 speech and Ronald Regan's 1981 speech as our anchors for a Wordfish model.

(c) Which of the documents is the most left wing? Which is the most right-wing? Are these results surprising? Why or why not?

(d) Re-create the "guitar plot" from Recitation 10. Describe the parameters estimated by Wordfish that lie on the axes of the plot.

(e) **Optional:** Estimate a linear regression with the Wordfish score as the dependent variable and binary variable indicating whether or not a President was a Democrat as an independent variable. Include a binary control variable for each president. If we use being a Democrat as a proxy for liberal ideology, how well did our Wordfish model do at capturing latent ideology?[2]

[1] `fit.stm` Represents the output of the STM model you fit in the preceding question.
[2] If it did well, then our proxy variable for ideology should be significant at at least a 5% level.

5. **Burstiness:** Here we evaluate the burstiness of several words using Arthur's corpus of treaties between Native American tribes and the U.S. Government. To evaluate burstiness we will use the `bursts` package and the user-written function `bursty` from Recitation 12 that visualizes the results.

   (a) Create a corpus from the treaties using the `readtext` command. For each of the words, "Seminole", "the", and "removal" use the `bursty` function to visualize the burst period(s) and levels. Also, for each of the plots include a brief interpretation about what the timing and level of the burst indicates about groups and events associated with the treaties. **Hint:** Look at the events and parties affected by the Indian Removal Act of 1830. You can use the following synopsis as a reference: Indian Treaties and the Removal Act of 1830